

What is Kubernetes?




- Kubernetes = popular container orchestrator
- Container Orchestration = Make many servers act like one
- Released by Google in 2014, maintained by large community
- Runs on top of Docker (usually) as a set of APIs in containers
- Provides API/CLI to manage containers across servers
- Many clouds provide it for you
- Many vendors make a "distribution" of it

Why Kubernetes?




- Review "Swarm Mode: Built-In Orchestration"
- Orchestration: Next logical step in journey to faster DevOps
- First, understand why you *may* need orchestration
- Not every solution needs orchestration
- Servers + Change Rate = Benefit of orchestration
- Then, decide which orchestrator
- If Kubernetes, decide which distribution
 - cloud or self-managed (Docker Enterprise, Rancher, OpenShift, Canonical, VMWare PKS)
 - Don't usually need pure upstream

Kubernetes or Swarm?



- Review "Swarm Mode: Built-In Orchestration"
- Kubernetes and Swarm are both container orchestrators
- Both are solid platforms with vendor backing
- Swarm: Easier to deploy/manage
- Kubernetes: More features and flexibility
- What's right for you? Understand both and know your requirements

Advantages of Swarm



- Comes with Docker, single vendor container platform
- Easiest orchestrator to deploy/manage yourself
- Follows 80/20 rule, 20% of features for 80% of use cases
- Runs anywhere Docker does:
 - local, cloud, datacenter
 - ARM, Windows, 32-bit
- Secure by default
- Easier to troubleshoot

Advantages of Kubernetes



- Clouds will deploy/manage Kubernetes for you
- Infrastructure vendors are making their own distributions
- Widest adoption and community
- Flexible: Covers widest set of use cases
- "Kubernetes first" vendor support
- "No one ever got fired for buying IBM"
 - Picking solutions isn't 100% rational
 - Trendy, will benefit your career
 - CIO/CTO Checkbox



Basic Terms: System Parts

- Kubernetes: The whole orchestration system
 - K8s "k-eights" or Kube for short
- Kubectl: CLI to configure Kubernetes and manage apps
 - Using "cube control" official pronunciation
- Node: Single server in the Kubernetes cluster
- Kubelet: Kubernetes agent running on nodes
- Control Plane: Set of containers that manage the cluster
 - Includes API server, scheduler, controller manager, etcd, and more
 - Sometimes called the "master"

Install Kubernetes Locally



- Kubernetes is a series of containers, CLI's, and configurations
- Many ways to install, lets focus on easiest for learning
- Docker Desktop: Enable in settings
 - Sets up everything inside Docker's existing Linux VM
- Docker Toolbox on Windows: MiniKube
 - Uses VirtualBox to make Linux VM
- Your Own Linux Host or VM: MicroK8s
 - Installs Kubernetes right on the OS

Kubernetes In A Browser



- Try <http://play-with-k8s.com> or katacoda.com in browser
- Easy to get started
- Doesn't keep your environment

Docker Desktop



- Runs/configures Kubernetes Master containers
- Manages kubectl install and certs
- Easily install, disable, and remove from Docker GUI

MiniKube



- Download Windows Installer from GitHub
- minikube-installer.exe
- minikube start
- Much like the docker-machine experience
- Creates a VirtualBox VM with Kubernetes master setup
- Doesn't install kubectl

MicroK8s



- Installs Kubernetes (without Docker Engine) on localhost (Linux)
- Uses snap (rather than apt or yum) for install
- Control the MicroK8s service via **microk8s.** commands
- kubectl accessible via **microk8s.kubectl**
- Add CoreDNS for services to work
 - **microk8s.enable dns**
- Add an alias to your shell (**.bash_profile**)
 - **alias kubectl=microk8s.kubectl**

Kubernetes Container Abstractions



- **Pod:** one or more containers running together on one Node
 - Basic unit of deployment. Containers are always in pods
- **Controller:** For creating/updating pods and other objects
 - Many types of Controllers inc. Deployment, ReplicaSet, StatefulSet, DaemonSet, Job, CronJob, etc.
- **Service:** network endpoint to connect to a pod
- **Namespace:** Filtered group of objects in cluster
- **Secrets, ConfigMaps, and more**

Kubernetes Run, Create, and Apply



- Kubernetes is evolving, and so is the CLI
- We get three ways to create pods from the kubectl CLI
 - > `kubectl run` (changing to be only for pod creation)
 - > `kubectl create` (create some resources via CLI or YAML)
 - > `kubectl apply` (create/update anything via YAML)
- For now we'll just use `run` or `create` CLI
- Later we'll learn YAML and pros/cons of each



Creating Pods with kubectl

- Are we working?
 - > `kubectl version`
- Two ways to deploy Pods (containers): Via commands, or via YAML
- Let's run a pod of the nginx web server!
 - > `kubectl create deployment my-nginx --image nginx`
- Let's list the pod
 - > `kubectl get pods`
- Let's see all objects
 - > `kubectl get all`

Pods -> ReplicaSet -> Deployment



Cleanup



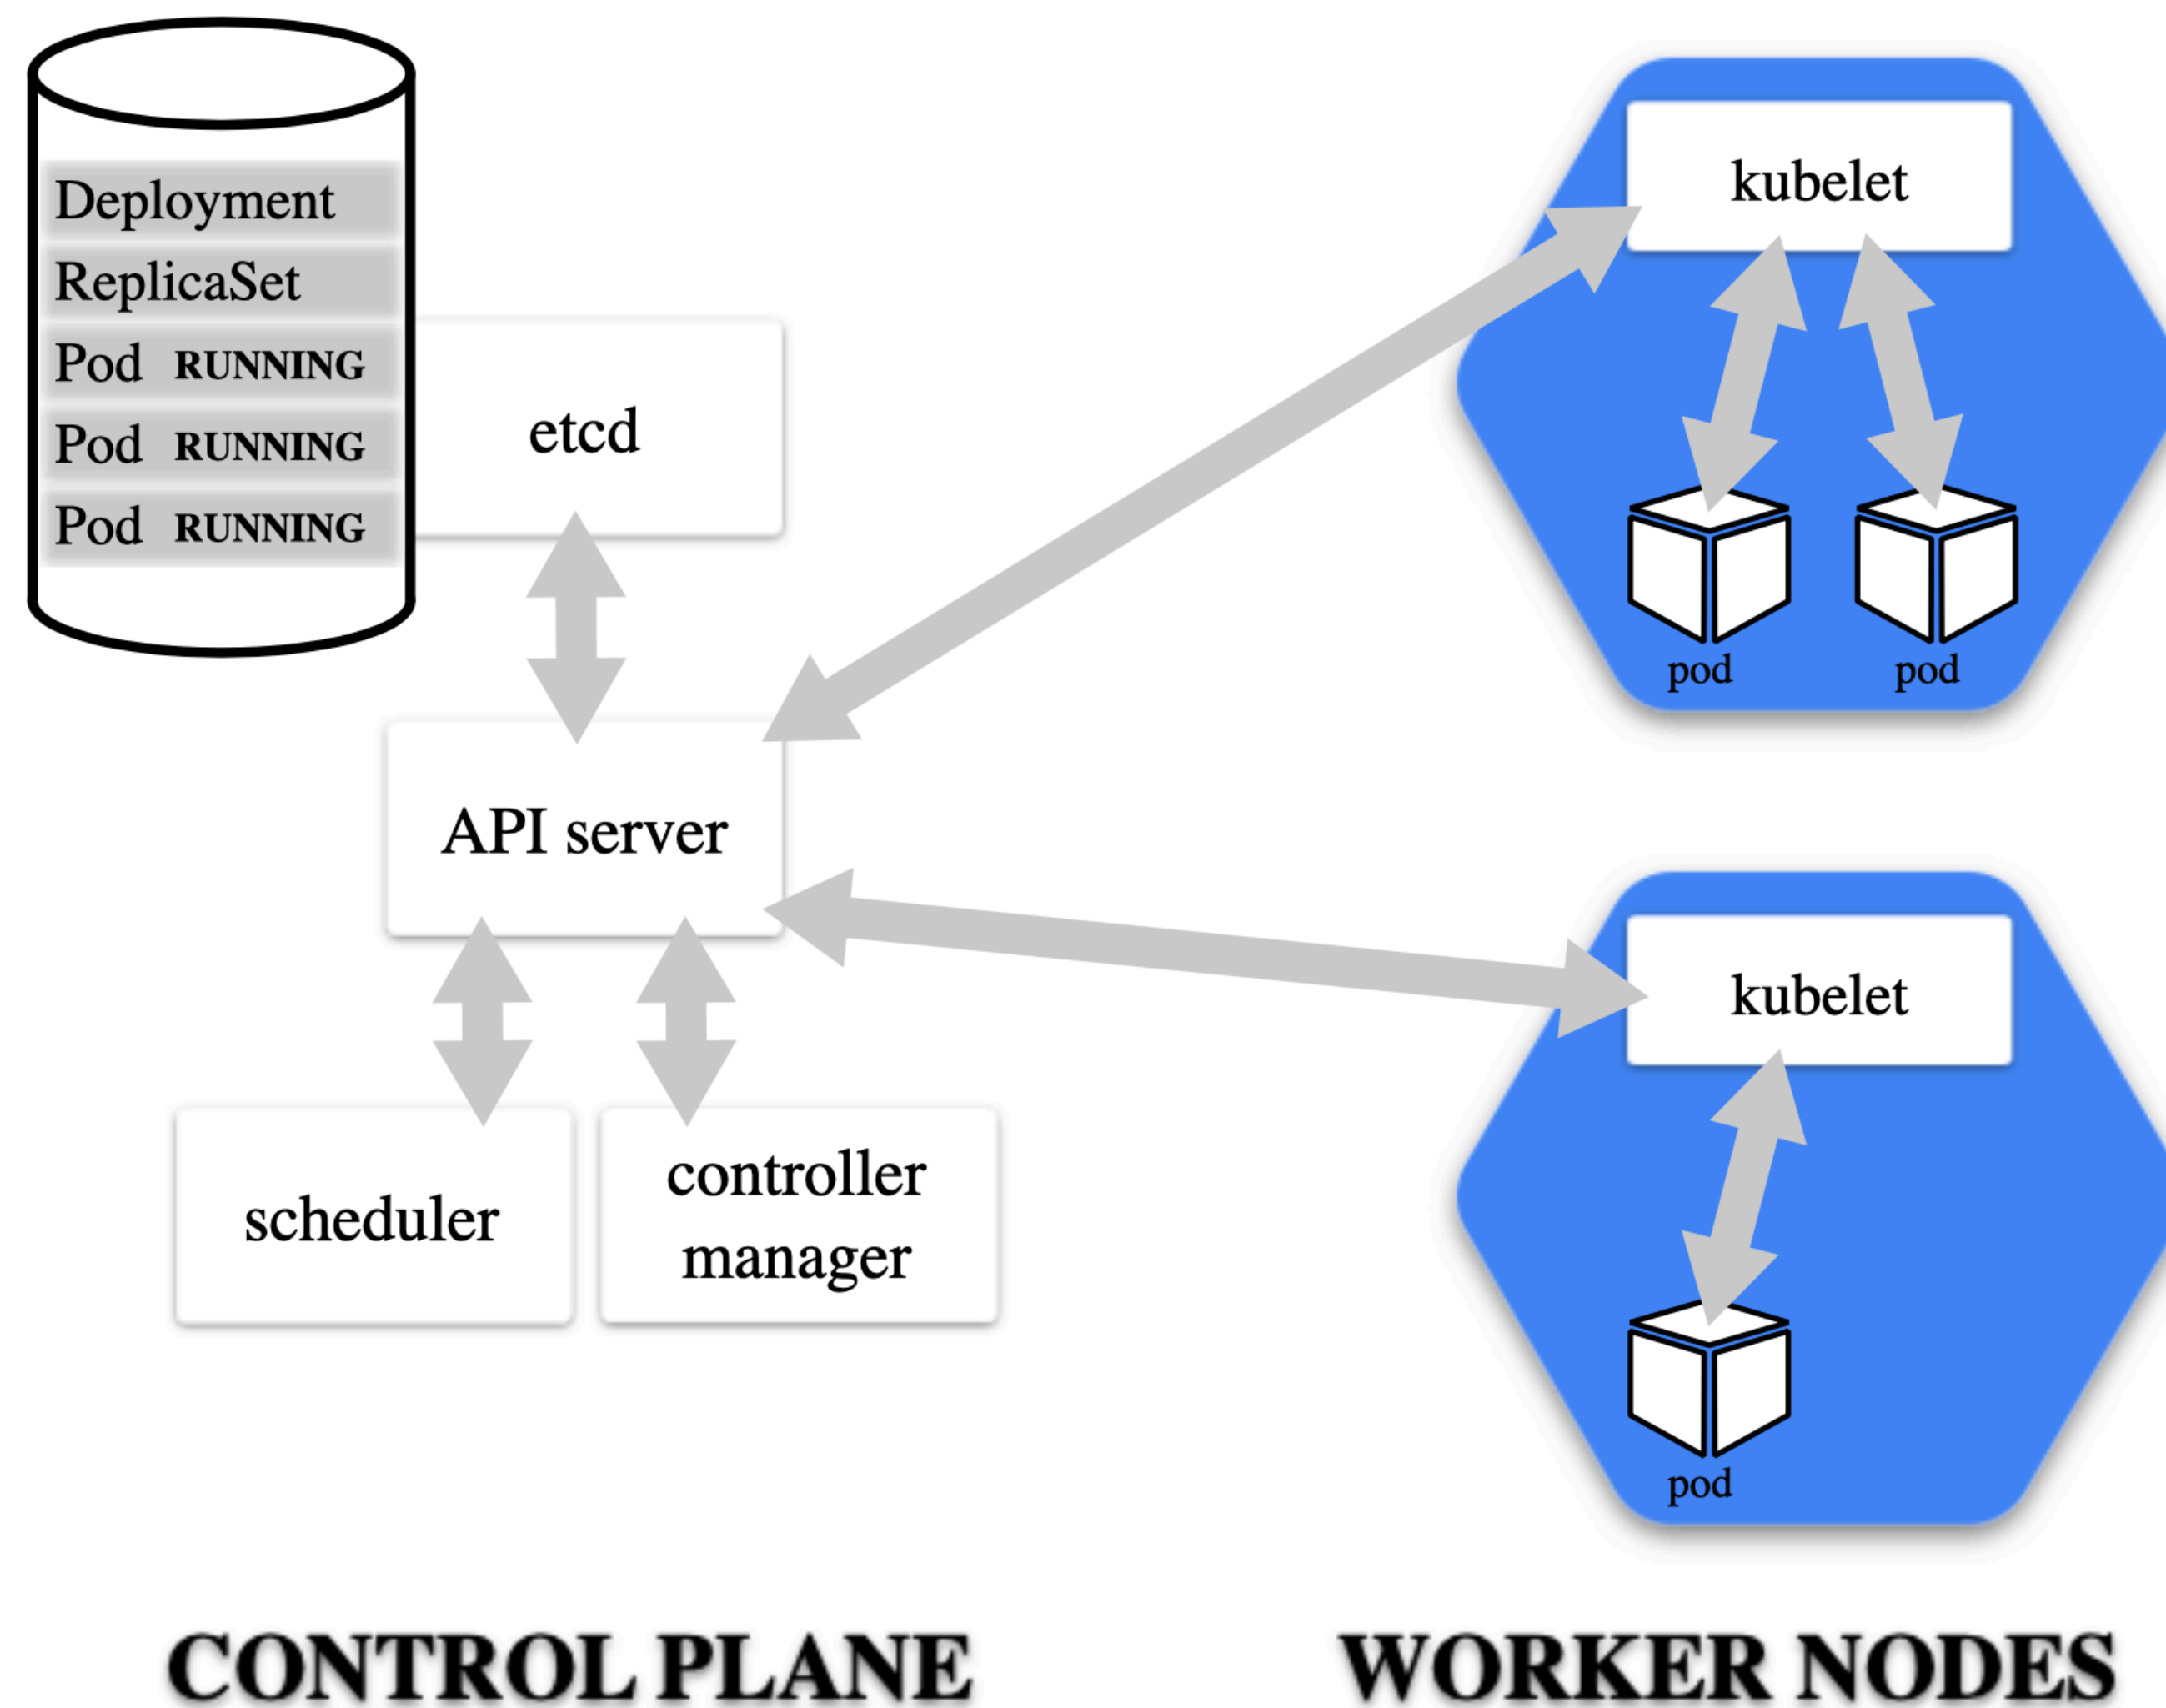
- Let's remove the Deployment
 - > `kubectl delete deployment my-nginx`

Scaling ReplicaSets



- Start a new deployment for one replica/pod
 - > `kubectl create deployment my-apache --image httpd`
- Let's scale it up with another pod
 - > `kubectl scale deploy/my-apache --replicas 2`
 - > `kubectl scale deployment my-apache --replicas 2`
- those are the same command
- `deploy = deployment = deployments`

What Just Happened? kubectl scale



Don't Cleanup



- We'll use these httpd containers in the next lecture



Inspecting Deployment Objects

- > `kubectl get pods`
- Get container logs
 - > `kubectl logs deployment/my-apache --follow --tail 1`
- Get a bunch of details about an object, including events!
 - > `kubectl describe pod/my-apache-xxxx-yyyy`
- Watch a command (without needing `watch`)
 - > `kubectl get pods -w`
- In a separate tab/window
 - > `kubectl delete pod/my-apache-xxxx-yyyy`
- Watch the pod get re-created

Cleanup



- Let's remove the Deployment
 - > `kubectl delete deployment my-apache`

Exposing Containers



- **kubectl expose** creates a **service** for existing pods
- A **service** is a stable address for pod(s)
- If we want to connect to pod(s), we need a **service**
- CoreDNS allows us to resolve **services** by name
- There are different types of **services**
 - ClusterIP
 - NodePort
 - LoadBalancer
 - ExternalName



Basic Service Types

- **ClusterIP** (default)
 - Single, internal virtual IP allocated
 - Only reachable from within cluster (nodes and pods)
 - Pods can reach service on apps port number
- **NodePort**
 - High port allocated on each node
 - Port is open on every node's IP
 - Anyone can connect (if they can reach node)
 - Other pods need to be updated to this port
- **These services are always available in Kubernetes**

More Service Types



- **LoadBalancer**
 - Controls a LB endpoint external to the cluster
 - Only available when infra provider gives you a LB (AWS ELB, etc)
 - Creates NodePort+ClusterIP services, tells LB to send to NodePort
- **ExternalName**
 - Adds CNAME DNS record to CoreDNS only
 - Not used for Pods, but for giving pods a DNS name to use for something outside Kubernetes
- **Kubernetes Ingress:** We'll learn later

Creating a ClusterIP Service



- Open two shell windows so we can watch this
 - > `kubectl get pods -w`
- In second window, lets start a simple http server using sample code
 - > `kubectl create deployment httpenv --image=bretfisher/httpenv`
- Scale it to 5 replicas
 - > `kubectl scale deployment/httpenv --replicas=5`
- Let's create a ClusterIP service (default)
 - > `kubectl expose deployment/httpenv --port 8888`

Inspecting ClusterIP Service



- Look up what IP was allocated
 - > `kubectl get service`
- Remember this IP is cluster internal only, how do we curl it?
- If you're on Docker Desktop (Host OS is not container OS)
 - > `kubectl run --generator=run-pod/v1 tmp-shell --rm -it --image bretfisher/netshoot -- bash`
 - > `curl httpenv:8888`
- If you're on Linux host
 - > `curl [ip of service]:8888`

Cleanup



- Leave the deployment there, we'll use it in the next Lecture

Create a NodePort Service



- Let's expose a NodePort so we can access it via the host IP (including localhost on Windows/Linux/macOS)
 - > `kubectl expose deployment/httpenv --port 8888 --name httpenv-np --type NodePort`
- Did you know that a NodePort service also creates a ClusterIP?
- These three service types are additive, each one creates the ones above it:
 - ClusterIP
 - NodePort
 - LoadBalancer

Add a LoadBalancer Service



- If you're on Docker Desktop, it provides a built-in LoadBalancer that publishes the `--port` on localhost
 - > `kubectl expose deployment/httpenv --port 8888 --name httpenv-lb --type LoadBalancer`
 - > `curl localhost:8888`
- If you're on kubeadm, minikube, or microk8s
 - No built-in LB
 - You can still run the command, it'll just stay at "pending" (but its NodePort works)

Cleanup



- Let's remove the Services and Deployment
 - > `kubectl delete service/httpenv service/httpenv-np`
 - > `kubectl delete service/httpenv-lb deployment/httpenv`



Kubernetes Services DNS

- Starting with 1.11, internal DNS is provided by CoreDNS
- Like Swarm, this is DNS-Based Service Discovery
- So far we've been using hostnames to access Services
 - > curl <hostname>
- But that only works for Services in the same Namespace
 - > kubectl get namespaces
- Services also have a FQDN
 - > curl <hostname>.<namespace>.svc.cluster.local

Assignment: Explore run get and logs



- Dry Run
 - > `kubectl create deployment nginx --image nginx --dry-run`
- Run does different things based on options
 - > `kubectl create deployment nginx --image nginx --dry-run --port 80 --expose`
- Only create a simple Pod, not a Deployment, ReplicaSet, etc.
 - > `kubectl run nginx-pod --generator=run-pod/v1 --image nginx`
- Get a shell in new Pod, remove on exit
 - > `kubectl run shell --generator=run-pod/v1 --rm -it --image busybox`

Assignment: Explore run get and logs



- Create a Deployment and ClusterIP Service in one line
 - > `kubectl run nginx2 --image nginx --replicas 2`
- Get multiple resources in one line
 - > `kubectl get deploy,pods`
- Get all pods, in wide format (gives more info)
 - > `kubectl get pods -o wide`
- Get all pods and show labels
 - > `kubectl get pods --show-labels`
-

Assignment: Explore run get and logs



- Better log viewing with stern
 - github.com/wercker/stern
 - > `kubectl run mydate --image bretfisher/date --replicas 3`
 - > `kubectl logs deployment/mydate`
 - > `stern mydate`

Cleanup



- Let's remove everything but the service/kubernetes
 - > kubectl get all
 - > kubectl delete deployment/nginx2 pod/nginx-pod

Run, Create, and Expose Generators



- These commands use helper templates called "generators"
- Every resource in Kubernetes has a specification or "spec"
 - > `kubectl create deployment sample --image nginx --dry-run -o yaml`
- You can output those templates with `--dry-run -o yaml`
- You can use those YAML defaults as a starting point
- Generators are "opinionated defaults"

Generator Examples



- Using dry-run with yaml output we can see the generators
 - > `kubectl create deployment test --image nginx --dry-run -o yaml`
 - > `kubectl create job test --image nginx --dry-run -o yaml`
 - > `kubectl expose deployment/test --port 80 --dry-run -o yaml`
 - You need the deployment to exist before this works

Cleanup



- Let's remove the Deployment
 - > `kubectl delete deployment test`



The Future of kubectl run

- Right now (1.12-1.15) run is in a state of flux
- The goal is to reduce its features to only create Pods
 - Right now it defaults to creating Deployments (with the warning)
 - It has lots of generators but they are all deprecated
 - The idea is to make it easy like **docker run** for one-off tasks
- It's not recommended for production
- Use for simple dev/test or troubleshooting pods

Old Run Confusion



- The generators activate different Controllers based on options
- Using dry-run we can see which generators are used
 - > `kubectl run test --image nginx --dry-run`
 - > `kubectl run test --image nginx --port 80 --expose --dry-run`
 - > `kubectl run test --image nginx --restart OnFailure --dry-run`
 - > `kubectl run test --image nginx --restart Never --dry-run`
 - > `kubectl run test --image nginx --schedule "*/1 * * * *" --dry-run`



Imperative vs. Declarative

- Imperative: Focus on *how* a program operates
- Declarative: Focus on *what* a program should accomplish
- Example: "I'd like a cup of coffee"
- Imperative: I boil water, scoop out 42 grams of medium-fine grounds, pour over 700 grams of water, etc.
- Declarative: "Barista, I'd like a a cup of coffee".
(Barista is the engine that works through the steps, including retrying to make a cup, and is only finished when I have a cup)

Kubernetes Imperative



- Examples: `kubectl run`, `kubectl create deployment`, `kubectl update`
 - We start with a state we know (no deployment exists)
 - We ask `kubectl run` to create a deployment
- Different commands are required to change that deployment
- Different commands are required per object
- Imperative is easier when you know the state
- Imperative is easier to get started
- Imperative is easier for humans at the CLI
- Imperative is NOT easy to automate

Kubernetes Declarative



- Example: `kubectl apply -f my-resources.yaml`
 - We don't know the current state
 - We only know what we want the end result to be (yaml contents)
- Same command each time (tiny exception for delete)
- Resources can be all in a file, or many files (apply a whole dir)
- Requires understanding the YAML keys and values
- More work than `kubectl run` for just starting a pod
- The easiest way to automate
- The eventual path to GitOps happiness

Three Management Approaches



- **Imperative commands:** `run`, `expose`, `scale`, `edit`, `create deployment`
 - Best for dev/learning/personal projects
 - Easy to learn, hardest to manage over time
- **Imperative objects:** `create -f file.yml`, `replace -f file.yml`, `delete...`
 - Good for prod of small environments, single file per command
 - Store your changes in git-based yaml files
 - Hard to automate
- **Declarative objects:** `apply -f file.yml` or `dir\`, `diff`
 - Best for prod, easier to automate
 - Harder to understand and predict changes

Three Management Approaches



- **Most Important Rule:**
 - Don't mix the three approaches
- **Bret's recommendations:**
 - Learn the Imperative CLI for easy control of local and test setups
 - Move to `apply -f file.yml` and `apply -f directory\` for prod
 - Store yaml in git, git commit each change before you apply
 - This trains you for later doing GitOps (where git commits are automatically applied to clusters)

kubectl apply



- Remember the three management approaches?
- Let's skip to full Declarative objects
- > kubectl apply -f filename.yml
- Why skip kubectl create, kubectl replace, kubectl edit?
- What I recommend ≠ all that's possible



Using kubectl apply

- create/update resources in a file
 - > `kubectl apply -f myfile.yaml`
- create/update a whole directory of yaml
 - > `kubectl apply -f myyaml/`
- create/update from a URL
 - > `kubectl apply -f https://bret.run/pod.yaml`
- Be careful, lets look at it first (browser or curl)
 - > `curl -L https://bret.run/pod`
 - Win PoSH? `start https://bret.run/pod.yaml`

Kubernetes Configuration YAML



- Kubernetes configuration file (YAML or JSON)
- Each file contains one or more manifests
- Each manifest describes an API object (deployment, job, secret)
- Each manifest needs four parts (root key:values in the file)

apiVersion:

kind:

metadata:

spec:



Building Your YAML Files

- **kind:** We can get a list of resources the cluster supports
 - > `kubectl api-resources`
- Notice some resources have multiple API's (old vs. new)
- **apiVersion:** We can get the API versions the cluster supports
 - > `kubectl api-versions`
- **metadata:** only name is required
- **spec:** Where all the action is at!



Building Your YAML spec

- We can get all the keys each kind supports
 - > `kubectl explain services --recursive`
- Oh boy! Let's slow down
 - > `kubectl explain services.spec`
- We can walk through the spec this way
 - > `kubectl explain services.spec.type`
- `spec:` can have sub `spec:` of other resources
 - > `kubectl explain deployment.spec.template.spec.volumes.nfs.server`
- We can also use docs
 - kubernetes.io/docs/reference/#api-reference



Dry Runs With Apply YAML

- New stuff, not out of beta yet (1.15)
- dry-run a create (client side only)
 - > `kubectl apply -f app.yml --dry-run`
- dry-run a create/update on server
 - > `kubectl apply -f app.yml --server-dry-run`
- see a diff visually
 - > `kubectl diff -f app.yml`



Labels and Annotations

- Labels goes under **metadata**: in your YAML
- Simple list of **key: value** for identifying your resource later by selecting, grouping, or filtering for it
- Common examples include **tier: frontend**, **app: api**, **env: prod**, **customer: acme.co**
- Not meant to hold complex, large, or non-identifying info, which is what **annotations** are for
- filter a get command
 - > **kubectl get pods -l app=nginx**
- apply only matching labels
 - > **kubectl apply -f myfile.yaml -l app=nginx**

Label Selectors



- The "glue" telling Services and Deployments which pods are theirs
- Many resources use Label Selectors to "link" resource dependencies
- You'll see these match up in the Service and Deployment YAML
- Use Labels and Selectors to control which pods go to which nodes
- Taints and Tolerations also control node placement

Cleanup



- Let's remove anything you created in this section
 - > `kubectl get all`
 - > `kubectl delete <resource type> / <resource name>`

Storage in Kubernetes



- Storage and stateful workloads are harder in all systems
- Containers make it both harder and easier than before
- **StatefulSets** is a new resource type, making Pods more sticky
- Bret's recommendation: avoid stateful workloads for first few deployments until you're good at the basics
 - Use db-as-a-service whenever you can



Volumes in Kubernetes

- Creating and connecting Volumes: 2 types
- **Volumes**
 - Tied to lifecycle of a Pod
 - All containers in a single Pod can share them
- **PersistentVolumes**
 - Created at the cluster level, outlives a Pod
 - Separates storage config from Pod using it
 - Multiple Pods can share them
- CSI plugins are the new way to connect to storage

Ingress



- None of our Service types work at OSI Layer 7 (HTTP)
- How do we route outside connections based on hostname or URL?
- Ingress Controllers (optional) do this with 3rd party proxies
- Nginx is popular, but Traefik, HAProxy, F5, Envoy, Istio, etc.
- Note this is still beta (in 1.15) and becoming popular
- Implementation is specific to Controller chosen

CRD's and The Operator Pattern



- You can add 3rd party Resources and Controllers
- This extends Kubernetes API and CLI
- A pattern is starting to emerge of using these together
- Operator: automate deployment and management of complex apps
- e.g. Databases, monitoring tools, backups, and custom ingresses

Higher Deployment Abstractions



- All our **kubectl** commands just talk to the Kubernetes API
- Kubernetes has limited built-in templating, versioning, tracking, and management of your apps
- There are now over 60 3rd party tools to do that, but many are defunct
- **Helm** is the most popular
- "**Compose on Kubernetes**" comes with Docker Desktop
- Remember these are optional, and your distro may have a preference
- Most distros support **Helm**

Templating YAML



- Many of the deployment tools have templating options
- You'll need a solution as the number of environments/apps grow
- Helm was the first "winner" in this space, but can be complex
- Official Kustomize feature works out-of-the-box (as of 1.14)
- **docker app** and compose-on-kubernetes are Docker's way

Kubernetes Dashboard



- Default GUI for "upstream" Kubernetes
 - github.com/kubernetes/dashboard
- Some distributions have their own GUI (Rancher, Docker Ent, OpenShift)
- Clouds don't have it by default
- Let's you view resources and upload YAML
- Safety first!

Kubectl Namespaces and Context



- Namespaces limit scope, aka "virtual clusters"
- Not related to Docker/Linux namespaces
- Won't need them in small clusters
- There are some built-in, to hide system stuff from **kubectl** "users"
 - > **kubectl get namespaces**
 - > **kubectl get all --all-namespaces**
- Context changes **kubectl** cluster and namespace
- See **~/.kube/config** file
 - > **kubectl config get-contexts**
 - > **kubectl config set***

Future of Kubernetes



- More focus on stability and security
 - 1.14, 1.15, largely dull releases (a good thing!)
 - Recent security audit has created backlog
- Clearing away deprecated features like kubectl run generators
- Improving features like server-side dry-run
- More and improved Operators
- Helm 3.0 (easier deployment, chart repos, libs)
- More declarative-style features
- Better Windows Server support
- More edge cases, kubeadm HA clusters

Related Projects



- Kubernetes has become the "differentencing and scheduling engine backbone" for so many new projects
- Knative - Serverless workloads on Kubernetes
- k3s - mini, simple Kubernetes
- k3OS - Minimal Linux OS for k3s
- Service Mesh - New layer in distributed app traffic for better control, security, and monitoring